

IMPROVING DEEP NEURAL NETWORKS: HYPERPARAMETER TUNING, REGULARIZATION AND OPTIMIZATION

By:
Satish Deshbhratar

BIAS / VARIANCE

- ❖ Applied ML: Iterative process to detect the hyperparameters
- ❖ Hyperparameters: learning rate, #iterations, #hidden layers/units...
- ❖ 100 → 1,000,000 : 60% training set – 20% development (cross validation) set – 20% test set
- ❖ 1,000,000+: 98% training set – 1% dev set – 1% test set
- ❖ PS: Mismatched dev / test distribution is bad. They have to come from the same distribution.
- ❖ PS: dev set is essential, but test set is not.
- ❖ High bias → Underfitting → High error in train → Train: 15% / Test: 16% (linear in a region)
- ❖ High variance → Overfitting → High error in difference → Train: 1% / Test: 11% (flexibility in a region)
- ❖ Bias variance tradeoff: They were used to be related inversely.
- ❖ High bias solutions: Bigger network / More layers / NN architecture / Different model / Longer train
- ❖ High variance solutions: More data / Regularization / NN architecture / Different model

REGULARIZATION

❖ *L2 Regularization*: $\frac{\lambda}{2m} ||w||_2^2 = \frac{\lambda}{2m} W^T \cdot W$

❖ *L1 Regularization*: $\frac{\lambda}{2m} ||w||_1$


❖ ***Frobinus Regularization***: $\frac{\lambda}{2m} ||W^{[l]}||_F^2 = \frac{\lambda}{2m} \sum_{i=1}^{n^{[l-1]}} \sum_{j=1}^{n^{[l]}} w_{ij}^{[l]2}$

❖ \rightarrow New backpropagation: $dW^{[l]} = \text{Old term} + \frac{\lambda}{m} w^{[l]}$: *Weight decay*

❖ Regularization vs overfitting: $\lambda \rightarrow +\infty, w^{[l]} \rightarrow 0, \Rightarrow$
Logistic regression (Single NN)


❖ Large $\lambda \rightarrow$ smoother decision boundry / smaller weights / remove over fitting

DROPOUT



- ❖ Concept: Remove some hidden units each iteration based on probability (inverted dropout)
- ❖ $d3 = \text{np.random.rand}(a3.\text{shape}[0], a3.\text{shape}[1]) < \text{keep_prob}$ #Generate 0's
- ❖ $a3 = a3 * d3 = \text{np.multiply}(a3, d3)$ #Element-wise multiplication
- ❖ $a3 = a3 / \text{keep_prob}$ #Keep the values on the same scale → Shrink weights
- ❖ PS: Each layer can have its own keep_prob; keep_prob=1 in first and last layer
- ❖ PS: Smaller keep_prob → More weights → Tend to overfit more
- ❖ PS: Dropout is only used in training

MINI-BATCH GRADIENT DESCENT



- ❖ Iterative model = Empiric model = Computational model
- ❖ Batch gradient descent: Mini-batch gradient descent with m #Too long per iteration
- ❖ Stochastic gradient descent: Mini-batch gradient descent with 1 #Lose vectorization speed, never converge
- ❖ $X \in (n_x, m)$; $Y \in (1, m)$; $X = [X^{(1)}, \dots, X^{(1000)}, X^{(1001)}, \dots, X^{(2000)} \dots]$

- ❖ for $t=1, \dots, \text{number of mini-batches}$: (each iteration, as if $m=1000$)
 - ❖ forward prop on $X^{[t]}$
 - ❖ compute cost $J^{\{t\}} = \frac{1}{1000} \dots$
 - ❖ backward prop

$X^{[1]} (n, 1000)$: First mini-batch

- ❖ Epoch = Pass through training set = Iteration

- ❖ PS: X and Y are similarly split

- ❖ PS: Bigger learning rate for stochastic \rightarrow Less noise

- ❖ $m < 2000 \rightarrow$ Batch gradient descent ; Typical size for number of batches $\rightarrow 2^6, 2^7, 2^8 \dots$ and fit in CPU/GPU


- ❖ Number of batches increase \rightarrow Noise decrease

1) Shuffle randomly

2) Partition

PS: If it is not divisible by the mini-batch size, the last one will have smaller values


```
permutation = list(np.random.permutation(m))
shuffled_X = X[:, permutation]
shuffled_Y = Y[:, permutation].reshape((1, m))
```



OTHER REGULARIZATION TECHNIQUES

❖ Data augmentation: Flip images horizontally, do random modifications, make distortions

❖ Early stopping:





❖ PS: $A = (A < 0,5)$ #Generate 0's and 1's

NORMALIZATION

❖ $\mu = \frac{1}{m} \sum_{i=1}^m X^{(i)}$; $X = X - \mu$ #Substract mean

❖ $\sigma^2 = \frac{1}{m} \sum_{i=1}^m X^{(i)2}$; $X = \frac{X}{\sigma^2}$ #Normalize variance

❖ PS: It must be the same for the train / test → Faster optimization (symmetric distribution)



INITIALIZATION


- ❖ Problem of Vanishing / Exploding gradients: Occurs when your derivatives are small/big
- ❖ Random initialization for W's to break symmetry; zeros for b's
- ❖ He initialization: $W^{[l]} = \text{np.random.randn(shape)} * \text{np.sqrt}(\frac{2}{n^{[l-1]}})$ #for RELU
- ❖ Another initialization: $W^{[l]} = \text{np.random.randn(shape)} * \text{np.sqrt}(\frac{1}{n^{[l-1]}})$ #for Tanh
- ❖ Xavier initialization: $W^{[l]} = \text{np.random.randn(shape)} * \text{np.sqrt}(\frac{1}{n^{[l-1]} + n^{[l]}})$

GRADIENT CHECKING

- ❖ Purpose: Find error in backpropagation implementation (gradient calculation) (slower)
- ❖ $w^{[1]}, b^{[1]}$.. \rightarrow Reshape in a big θ vector
- ❖ $dw^{[1]}, db^{[1]}$.. \rightarrow Reshape in a big $d\theta$ vector
- ❖ $J(w^{[1]}, b^{[1]}, dw^{[1]}, db^{[1]}...)$ $\rightarrow J(\theta)$
- ❖ for i: $d\theta_{approx}(i) = \frac{J(\theta_1, \dots, \theta_i + \epsilon, \dots) - J(\theta_1, \dots, \theta_i - \epsilon, \dots)}{2 \epsilon} = d\theta(i) = \frac{\partial J}{\partial \theta_i}$
- ❖ Check: $\frac{\|d\theta_{approx} - d\theta\|_2}{\|d\theta_{approx}\|_2 + \|d\theta\|_2} \leq \epsilon = 10^{-7}$
- ❖ PS: Grad check only in debug (not in train); Run grad check without dropout
- ❖ If grad check fails, look at components $db^{[l]}$ and $dw^{[l]}$
- ❖ $\sum_k \sum_j w_{j,k}^{[l]2} \Leftrightarrow \text{np.sum(np.square(w))}$; $\|x\|_2 \Leftrightarrow \text{np.linalg.norm(x)}$

EXPONENTIALLY WEIGHTED AVERAGE

- ❖ $V_0 = 0; V_t = \beta V_{t-1} + (1 - \beta)\theta_t$: *Averaging over last $\frac{1}{1-\beta}$ days' temperature*
- ❖ $\beta \nearrow \rightarrow$ Average graph becomes smoother and shift slightly to the right
- ❖ Bias correction: $\frac{V_t}{1-\beta^t}$ because V_t will be far from θ_1 at first



GRADIENT DESCENT WITH MOMENTUM


❖ For t:

❖ Compute dw , db on current mini-batch

$$❖ v_{dw} = \beta v_{dw} + (1 - \beta) v_{dw}$$

$$❖ v_{db} = \beta v_{db} + (1 - \beta) v_{db}$$

$$❖ w = w - \alpha v_{dw}; b = b - \alpha v_{db}$$



❖ We will be averaging on the vertical (around 0) and on the horizontal (straight forward)

❖ PS: Bias correction isn't needed because after few iterations we will be okay

❖ PS: $\beta=0,9$ in practice

❖ Another formulation:

$$❖ v_{dw} = \beta v_{dw} + v_{dw}$$

$$❖ w = w - \frac{\alpha}{1-\beta} v_{dw}$$

❖ PS: Momentum can be applied with any GD method

RMSPROP

❖ RMSprop = Root Mean Squared Propagation (Same objective as momentum)


❖ For t:

❖ Compute dw , db on current mini-batch

❖ $s_{dw} = \beta s_{dw} + (1 - \beta)dw^2$ #Element-wise power

❖ $s_{db} = \beta s_{db} + (1 - \beta)db^2$

❖ $w = w - \alpha \frac{dw}{\sqrt{s_{dw} + \epsilon}}$; $b = b - \alpha \frac{db}{\sqrt{s_{db} + \epsilon}}$ #We can get a bigger α now; ϵ to avoid divergence



ADAM

❖ ADAM = Momentum + RMSprop

❖ α needs to be tuned

❖ $\beta_1 = 0,9$ (Momentum)

❖ $\beta_2 = 0,999$ (RMSprop)

❖ $\epsilon = 10^{-8}$ (Doesn't matter much)

```
vdW = 0, vdb = 0
sdW = 0, sdb = 0
on iteration t:
    # can be mini-batch or batch gradient descent
    compute dw, db on current mini-batch

    vdW = (beta1 * vdW) + (1 - beta1) * dw      # momentum
    vdb = (beta1 * vdb) + (1 - beta1) * db      # momentum

    sdW = (beta2 * sdW) + (1 - beta2) * dw^2    # RMSprop
    sdb = (beta2 * sdb) + (1 - beta2) * db^2    # RMSprop

    vdW = vdW / (1 - beta1^t)                  # fixing bias
    vdb = vdb / (1 - beta1^t)                  # fixing bias

    sdW = sdW / (1 - beta2^t)                  # fixing bias
    sdb = sdb / (1 - beta2^t)                  # fixing bias
```

LEARNING RATE DECAY

❖ Concept: At first, you can have a bigger value of α and in the end you can make it smaller

$$\alpha = \frac{1}{1 + \text{decay rate} * \text{epoch number}} \alpha_0$$

$$\alpha = 0,95^{\text{epoch number}} \alpha_0$$


$$\alpha = \frac{k}{\sqrt{\text{epoch number}}} \alpha_0$$

❑ We can use manual decay too

❑ It is not evident to find a local optimum (derivate = 0) in a high dimensional space

❑ Saddle points are not really a problem

❑ Plateau: Region where the derivate is close to 0 for a long time (slows the algorithm)



HYPERPARAMETERS

❖ PS: We try 25 couples of hyperparameters

❖ 1st priority: α

❖ 2nd priority: β momentum, #hidden units, mini-batch size

❖ 3rd priority: #layers, learning rate decay

❖ Course to fine search process: If a couple of hyperparameters is good, we zoom in their region and pick more random values

❖ Appropriate scale for hyperparameters: α for example goes from 0,0001 to 1, we choose a logarithmic scale for a better distribution: $r = (\log a - \log b) * \text{np.random.rand}() + \log b \rightarrow \alpha = 10^r$

❖ $\beta = 0,9000 \rightarrow 0,9005 \sim$ Generate 10 values



❖ $\beta = 0,9990 \rightarrow 0,9995 \sim$ Generate 1000 values

❖ PS: Another method is: $r = \text{np.random.rand}(\text{interval})$

➔ Linear scale is bad

❖ Hyperparameters approach: Panda (Train one model at a time) vs Caviar (Train many models in parallel)

❖ PS: We use Caviar approach if we have high computational resources



BATCH NORMALIZATION

❖ For $z^{(i)}$:

❖ $\mu = \frac{1}{m} \sum_{i=1}^m Z^{(i)} ; \sigma^2 = \frac{1}{m} \sum_{i=1}^m Z^{(i)^2} ;$

❖ $Z_{norm}^{(i)} = \frac{Z^{(i)} - \mu}{\sqrt{\sigma^2 + \epsilon}}$ # ϵ to avoid division by zero

❖ $\tilde{Z}^{(i)} = \gamma Z_{norm}^{(i)} + \beta$ # To manipulate the mean and the variance # γ and β are learnable parameters

❖ Use $\tilde{Z}^{(i)}$ instead of $Z^{(i)}$ to make inputs belong to other distribution

❑ We use batch normalization in mini-batches

❑ The parameter b becomes 0, since $\text{mean}(Z)=0 \rightarrow$ No need for b parameter

❑ Shape of β and γ is $(n^{[l]}, 1)$

❑ We can use Momentum, RMSprop and ADAM with Batch Normalization



❑ Covariate shift problem: Training your algorithm on black cats \rightarrow It will be bad with colored cats

❑ \rightarrow BN reduces this problem by guaranteeing that they will have mean 0 and variance 1: Stability

❑ BN has a slight regularization effects (not recommended). BN adds a slight noise.

❑ In test, we use BN of all the layers in the training using exponentially weighted average

MULTICLASS CLASSIFICATION — SOFTMAX REGRESSION

- ❖ $C = \# \text{Classes} \{0, 1, 3, 4\} \rightarrow$ Output layer will have C units ($n^{[L]} = C$)
- ❖ \rightarrow Last layer = Output layer = Softmax layer:
$$\begin{bmatrix} p(C = 0/X) \\ p(C = 1/X) \\ p(C = 2/X) \\ p(C = 3/X) \end{bmatrix}; \sum_{i=1}^C p(C = i/X) = 1$$
- ❖ Activation function = Softmax activation function:
$$a_i^{[L]} = \frac{e^{z_i^{[L]}}}{\sum_{j=1}^C (e^{z_j^{[L]}})}$$
- ❖ PS: Usually activation functions take floats and return floats, here operation is on vectors
- ❖ Softmax $\begin{matrix} 0,8 \\ 0,1 \\ 0,05 \\ 0,05 \end{matrix}$ vs Hardmax $\begin{matrix} 1 \\ 0 \\ 0 \\ 0 \end{matrix} \rightarrow$ Softmax is generalization of sigmoid with $C \neq 2$
- ❖ New loss function: Likelihood function: $L(\hat{y}, y) = - \sum_{j=1}^C y_j \log \hat{y}_j$
- ❖ New backpropagation (gradient descent): $dZ^{[L]} = \hat{y} - y$

TENSORFLOW (1)

- ❖ **Framework** = Library that contain DL functions (Caffe, Kera, Tensorflow)
- ❖ **Choice**: ease of programming (dev and deployment), running speed, truly open (open source + good governance{will stay open source})
- ❖ `import tensorflow as tf`
- ❖ `w=tf.Variable(value,dtype=tf.float32)` #Define a parameter to optimize
- ❖ `cost_function = #Function of w here`
- ❖ `train=tf.train.GradientDescentOptimizer(learning_rate).minimize(cost_function)` #Define learning algo
 - ❖ `init=tf.global_variables_initializer()`
 - ❖ `session=tf.Session()` #Start a tf session
 - ❖ `session.run(init)` #Initialize variables
 - ❖ `session.run(w)` #Run session to calculate w
- ❖ `for i in range(numberOfIteration): session.run(train)` #Calculate grad desc

Idiomatic lines (Always in your code)

TENSORFLOW (2)

- ❖ PS: We prepare forward for Tensorflow, and it calculates backward for us
- ❖ `constant=tf.constant(value)`
- ❖ `x=tf.placeholder(tf.float32,size)` #A variable to assign value to, later
- ❖ `→ session.run(train,feed_dict={x:valuesToAssign,y:v})` #Get data in cost_function i.e. in train
- ❖ PS: float 5. \neq int 5 (in Python)
- ❖ `with tf.Session as session:`
 - ❖ `session.run(init)`
 - ❖ `Print(session.run(w))`
- ❖ Computation graph: What happens in tensorflow: It can compute backprop thanks to graph
- ❖ PS: In tf, we can't print(a+b) just like that, we have to run a session to be able to print it
- ❖ `tf.one_hot(Y,numberOfClasses,axis=0)` #Turn Y to a classification matrix

Another way to write the idiomatic lines

TENSORFLOW (3)

- ❖ `w=tf.get_variable(shape,initializer=tf.typeOfInitialization())`
- ❖ PS: `[n_x,None]` dimensions \Leftrightarrow number of lines = `n_x`, number of colons = `#idc`
- ❖ `tf.matmul(A,B)` `#A*B`
- ❖ `tf.nn.activationFunction(A)` `#Apply activationFunction on A`
- ❖ `tf.nn.softmax_cross_entropy_with_logits(logits=y,labels=z)` `#Compute cost`
- ❖ `_, a=#Something that return two outputs here` `#First return = #idc`
- ❖ PS: Main classes in tensorflow: Tensors (variables) and operators (functions/methods)